

# POWERAPI, Deploying Software-defined Power Meters for the Cloud

---

Open-source for the Cloud – ESIEA, Paris

23<sup>rd</sup> May, 2017 – 11:00

Maxime COLMANT  
Romain ROUYOY  
Lionel SEINTURIER

UNIVERSITY LILLE 1 / INRIA  
UNIVERSITY LILLE 1 / INRIA / IUF  
UNIVERSITY LILLE 1 / INRIA / IUF


# TABLE OF CONTENTS

1. Introduction
2. Contributions
3. Conclusion

# INTRODUCTION

---

# THE GLOBAL ICT<sup>1</sup> FOOTPRINT<sup>2</sup>



<sup>1</sup>Information and Communications Technology

<sup>2</sup>The Climate Group. *SMART 2020: Enabling the low carbon economy in the information age*. 2008.

# MULTI-CORE CPU ARCHITECTURES ARE EVERYWHERE!



Smartphone




Laptop



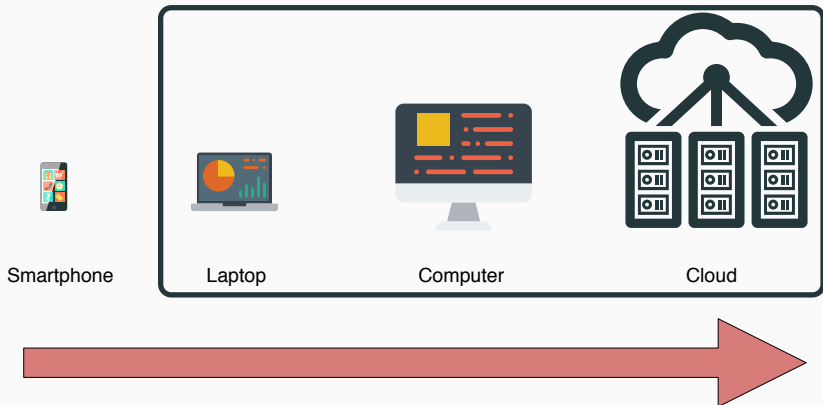
Computer



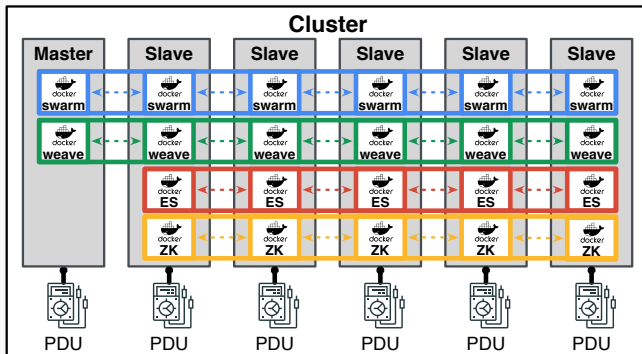
Cloud




# MULTI-CORE CPU ARCHITECTURES ARE EVERYWHERE!



# CASE STUDY




# CASE STUDY







# WHAT IS A MULTI-CORE SOFTWARE SYSTEM?




# WHAT IS A MULTI-CORE SOFTWARE SYSTEM?



# WHAT IS A MULTI-CORE SOFTWARE SYSTEM?




# WHAT IS A MULTI-CORE SOFTWARE SYSTEM?




# RESEARCH QUESTIONS

**RQ1:** Can we model the software power consumption regardless of the underlying architecture?



# RESEARCH QUESTIONS


**RQ2:** Can we propose a uniform view of the service power consumption?



## CONTRIBUTIONS


---

**RQ1:** Can we model the software power consumption regardless of the underlying architecture?





**RQ1:** Can we model the software power consumption regardless of the underlying architecture?



## Learning CPU Power Models

# WHAT IS A POWER MODEL?

- Math. function (metrics)  $\Rightarrow$  Power

# WHAT IS A POWER MODEL?

- Math. function (metrics)  $\Rightarrow$  Power
- Mostly linear

Univariate:  $P = a_x + b$

Multivariate:  $P = a_x + b_y + c$

# WHAT IS A POWER MODEL?

- Math. function (metrics)  $\Rightarrow$  Power
- Mostly linear

$$\text{Univariate: } P = a_x + b$$

$$\text{Multivariate: } P = a_x + b_y + c$$

- Or polynomial

$$P = a_{x^2} + b_x + c$$

# WHAT IS A POWER MODEL?

- Math. function (metrics)  $\Rightarrow$  Power

- Mostly linear

$$\text{Univariate: } P = a_x + b$$

$$\text{Multivariate: } P = a_x + b_y + c$$

- Or polynomial

$$P = a_{x^2} + b_x + c$$

- CPU metrics

From HW sensors (motherboard, power meters)

From Hardware Performance Counters (HPCs)

# WHAT IS A POWER MODEL?

- Math. function (metrics)  $\Rightarrow$  Power
- Mostly linear

$$\text{Univariate: } P = a_x + b$$

$$\text{Multivariate: } P = a_x + b_y + c$$

- Or polynomial

$$P = a_{x^2} + b_x + c$$

- CPU metrics

From HW sensors (motherboard, power meters)

From Hardware Performance Counters (HPCs)

- [Nou14]<sup>3</sup>:  $P_{cpu}^{app} = 0.7 * TDP * CPU_{stats}$

---

<sup>3</sup>A. Nouredine. “Towards a Better Understanding of the Energy Consumption of Software Systems”. PhD thesis. Université des Sciences et Technologie de Lille - Lille I, 2014.

# EXISTING CPU POWER MODELS

Ref.	Processor(s)	Feature(s)	Regression(s)	Benchmarks
[Ber+10]	Core 2 Duo	14 PCs regrouped by component	multiple linear by component	sampl.: $\mu$ -benchs eval.: SPEC CPU 06
[Col+15]	Xeon W3520 & i3 2120	non-halted cycles reference cycles	polynomial	sampl.: stress eval.: PARSEC, SPECjbb
[CM05]	XScale PXA255	5 PCs	multiple linear	eval.: SPEC CPU 00, Java CDC/CLDC
[Dol+15]	Xeon E3-1275	3 PCs HW sensors	linear	sampl.: linpack, stream, iperf, IOR eval.: Quantum Espresso
[ERK06]	Turion, Itanium 2	HW sensors	multiple linear	sampl.: Gamut eval.: SPECS, Matrix, Stream
[IM03]	Pentium 4	15 PCs	multiple linear	eval.: $\mu$ -benchs, AbiWord, Mozilla, Gnumeric
[RRK08]	Core 2 Duo & Xeon, Itanium 2, Turion	HW sensors PCs	multiple linear	sampl.: calibration suite eval.: SPECS, stream, Nsort
[Yan+14]	Xeon E5620 & E7530	7 components 91 preselected	support vector	sampl.: NPB, IOzone, CacheBench eval.: SPEC CPU 06, IOzone
[Zha+14]	Sandy Bridge	non-halted cycles	linear	eval.: Google, SPEC CPU 06
???	ARM	???	???	???

Only for Intel or AMD architectures

# EXISTING CPU POWER MODELS

Ref.	Processor(s)	Feature(s)	Regression(s)	Benchmarks
[Ber+10]	Core 2 Duo	14 HPCs regrouped by component	multiple linear by component	<i>sampl.</i> : $\mu$ -benchs <i>eval.</i> : SPEC CPU 06
[Col+15]	Xeon W3520 & i3 2120	non-halted cycles reference cycles	polynomial	<i>sampl.</i> : stress <i>eval.</i> : PARSEC, SPECjbb
[CM05]	XScale PXA255	5 HPCs	multiple linear	<i>eval.</i> : SPEC CPU 00, Java CDC/CLDC
[Dol+15]	Xeon E3-1275	3 HPCs HW sensors	linear	<i>sampl.</i> : linpack, stream, iperf, IOR <i>eval.</i> : Quantum Espresso
[ERK06]	Turion, Itanium 2	HW sensors	multiple linear	<i>sampl.</i> : Gamut <i>eval.</i> : SPECS, Matrix, Stream
[IM03]	Pentium 4	15 HPCs	multiple linear	<i>eval.</i> : $\mu$ -benchs, AbiWord, Mozilla, Gnumeric
[RRK08]	Core 2 Duo & Xeon, Itanium 2, Turion	HW sensors HPCs	multiple linear	<i>sampl.</i> : calibration suite <i>eval.</i> : SPECS, stream, Nsort
[Yan+14]	Xeon E5620 & E7530	7 components 91 preselected	support vector	<i>sampl.</i> : NPB, IOzone, CacheBench <i>eval.</i> : SPEC CPU 06, IOzone
[Zha+14]	Sandy Bridge	non-halted cycles	linear	<i>eval.</i> : Google, SPEC CPU 06

**HW sensors: coarse-grained CPU metrics**



# EXISTING CPU POWER MODELS

Ref.	Processor(s)	Feature(s)	Regression(s)	Benchmarks
[Ber+10]	Core 2 Duo	14 HPCs regrouped by component	multiple linear by component	<i>sampl.</i> : $\mu$ -benchs <i>eval.</i> : SPEC CPU 06
[Col+15]	Xeon W3520 & i3 2120	non-halted cycles reference cycles	polynomial	<i>sampl.</i> : stress <i>eval.</i> : PARSEC, SPECjbb
[CM05]	XScale PXA255	5 HPCs	multiple linear	<i>eval.</i> : SPEC CPU 00, Java CDC/CLDC
[Dol+15]	Xeon E3-1275	3 HPCs HW sensors	linear	<i>sampl.</i> : linpack, stream, iperf, IOR <i>eval.</i> : Quantum Espresso
[ERK06]	Turion, Itanium 2	HW sensors	multiple linear	<i>sampl.</i> : Gamut <i>eval.</i> : SPECS, Matrix, Stream
[IM03]	Pentium 4	15 HPCs	multiple linear	<i>eval.</i> : $\mu$ -benchs, AbiWord, Mozilla, Gnumeric
[RRK08]	Core 2 Duo & Xeon, Itanium 2, Turion	HW sensors HPCs	multiple linear	<i>sampl.</i> : calibration suite <i>eval.</i> : SPECS, stream, Nsort
[Yan+14]	Xeon E5620 & E7530	7 components 91 preselected	support vector	<i>sampl.</i> : NPB, IOzone, CacheBench <i>eval.</i> : SPEC CPU 06, IOzone
[Zha+14]	Sandy Bridge	non-halted cycles	linear	<i>eval.</i> : Google, SPEC CPU 06

**HPCs: fine-grained CPU metrics**

# EXISTING CPU POWER MODELS

Ref.	Processor(s)	Feature(s)	Regression(s)	Benchmarks
[Ber+10]	Core 2 Duo	14 HPCs regrouped by component	multiple linear by component	<i>sampl.</i> : $\mu$ -benchs <i>eval.</i> : SPEC CPU 06
[Col+15]	Xeon W3520 & i3 2120	non-halted cycles reference cycles	polynomial	<i>sampl.</i> : stress <i>eval.</i> : PARSEC, SPECjbb
[CM05]	XScale PXA255	5 HPCs	multiple linear	<i>eval.</i> : SPEC CPU 00, Java CDC/CLDC
[Dol+15]	Xeon E3-1275	3 HPCs HW sensors	linear	<i>sampl.</i> : linpack, stream, iperf, IOR <i>eval.</i> : Quantum Espresso
[ERK06]	Turion, Itanium 2	HW sensors	multiple linear	<i>sampl.</i> : Gamut <i>eval.</i> : SPECS, Matrix, Stream
[IM03]	Pentium 4	15 HPCs	multiple linear	<i>eval.</i> : $\mu$ -benchs, AbiWord, Mozilla, Gnumeric
[RRK08]	Core 2 Duo & Xeon, Itanium 2, Turion	HW sensors HPCs	multiple linear	<i>sampl.</i> : calibration suite <i>eval.</i> : SPECS, stream, Nsort
[Yan+14]	Xeon E5620 & E7530	7 components 91 preselected	support vector	<i>sampl.</i> : NPB, IOzone, CacheBench <i>eval.</i> : SPEC CPU 06, IOzone
[Zha+14]	Sandy Bridge	non-halted cycles	linear	<i>eval.</i> : Google, SPEC CPU 06

**Power models are mostly linear**

# EXISTING CPU POWER MODELS

Ref.	Processor(s)	Feature(s)	Regression(s)	Benchmarks
[Ber+10]	Core 2 Duo	14 HPCs regrouped by component	multiple linear by component	sampl.: $\mu$ -benchs eval.: <b>SPEC CPU 06</b>
[Col+15]	Xeon W3520 & i3 2120	non-halted cycles reference cycles	polynomial	sampl.: stress eval.: PARSEC, <b>SPECjbb</b>
[CM05]	XScale PXA255	5 HPCs	multiple linear	eval.: <b>SPEC CPU 00</b> , Java CDC/CLDC
[Dol+15]	Xeon E3-1275	3 HPCs HW sensors	linear	sampl.: linpack, stream, iperf, IOR eval.: Quantum Espresso
[ERK06]	Turion, Itanium 2	HW sensors	multiple linear	sampl.: Gamut eval.: <b>SPECs</b> , Matrix, Stream
[IM03]	Pentium 4	15 HPCs	multiple linear	eval.: $\mu$ -benchs, AbiWord, Mozilla, Gnumeric
[RRK08]	Core 2 Duo & Xeon, Itanium 2, Turion	HW sensors HPCs	multiple linear	sampl.: calibration suite eval.: <b>SPECs</b> , stream, Nsort
[Yan+14]	Xeon E5620 & E7530	7 components 91 preselected	support vector	sampl.: NPB, IOzone, CacheBench eval.: <b>SPEC CPU 06</b> , IOzone
[Zha+14]	Sandy Bridge	non-halted cycles	linear	eval.: <b>Google</b> , <b>SPEC CPU 06</b>

Non free or private workloads

## 1. Portability

1. Portability
2. Accuracy

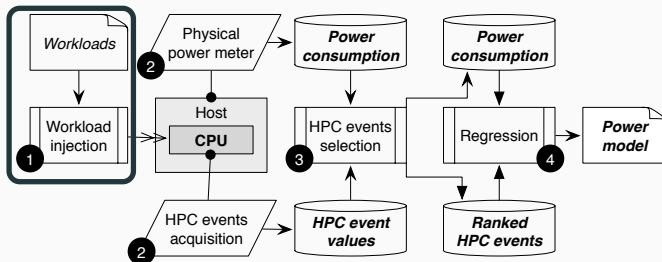
1. Portability
2. Accuracy
3. Reproducibility

1. Portability
2. Accuracy
3. Reproducibility

Towards an automatic approach for learning CPU power models

# OUR APPROACH:

## OPEN-TESTBED TO AUTOMATICALLY LEARN POWER MODELS



### 1 Input workload injection

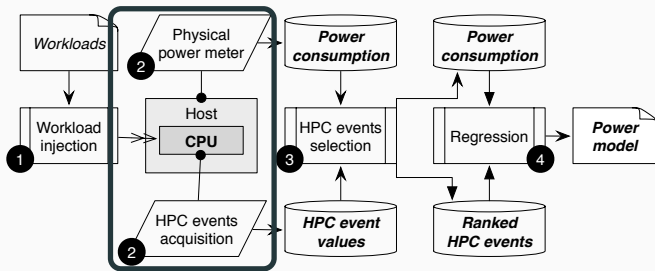
- Configurable
- PARSEC (open-source, multi-threaded)<sup>4</sup>
- Run several applications (x264, vips, etc.)

<sup>4</sup>C. Bienia et al. "PARSEC 2.0: A New Benchmark Suite for Chip-Multiprocessors". In: *Proceedings of the 5th Annual Workshop on Modeling, Benchmarking and Simulation*. 2009.



# OUR APPROACH:

## OPEN-TESTBED TO AUTOMATICALLY LEARN POWER MODELS



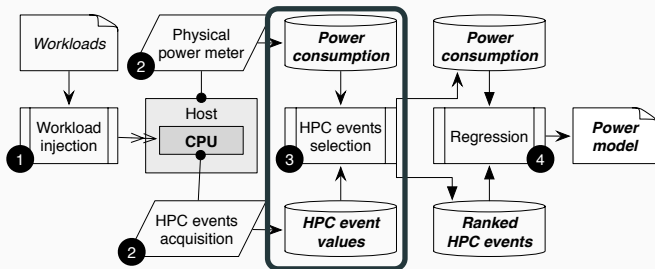
### ② Acquisition of raw input metrics

- Automatically explore the high number of the available HPCs (Xeon W3520: 514 HPCs)
- Take care of HPC multiplexing<sup>5</sup>

<sup>5</sup>Intel. *Intel 64 and IA-32 Architectures Software Developer's Manual*. 2015.

# OUR APPROACH:

## OPEN-TESTBED TO AUTOMATICALLY LEARN POWER MODELS




### 3 Selection of relevant HPCs

- Pearson coefficient ( $\text{HPC} \Leftrightarrow \text{Power}$ )
- 1<sup>st</sup> phase: quickly filtering out uncorrelated HPCs ( $< 0.5$ )  
(Xeon W3250: 253 left out)
- 2<sup>nd</sup> phase: full sampling for the remaining HPCs

# OUR APPROACH:

## OPEN-TESTBED TO AUTOMATICALLY LEARN POWER MODELS


Pearson coefficients of the Top-30 correlated events for the PARSEC benchmarks on a Xeon W3520.



# OUR APPROACH:

## OPEN-TESTBED TO AUTOMATICALLY LEARN POWER MODELS


Pearson coefficients of the Top-30 correlated events for the PARSEC benchmarks on a Xeon W3520.





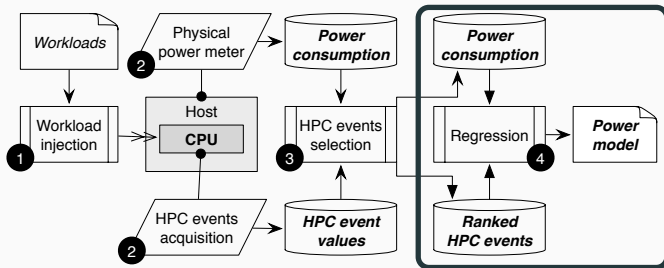
# OPEN-TESTBED TO AUTOMATICALLY LEARN POWER MODELS

Pearson coefficients of the Top-30 correlated events for the PARSEC benchmarks on a Xeon W3520.



# OUR APPROACH:

## OPEN-TESTBED TO AUTOMATICALLY LEARN POWER MODELS



### ④ Power model inference


- Minimize the number of HPCs
- Robust ridge regression (SotA?)

# OUR APPROACH:

## OPEN-TESTBED TO AUTOMATICALLY LEARN POWER MODELS

Average error per combination of HPCs for **freqmine**, **fluidanimate**, **facesim** on a Xeon W3520.

$$P_{idle} = 92\text{ W} ; P_{CPU} = \frac{1.40 \cdot \text{HPC (l1i:reads)}}{10^8} + \frac{7.29 \cdot \text{HPC (lsd:inactive)}}{10^9}$$

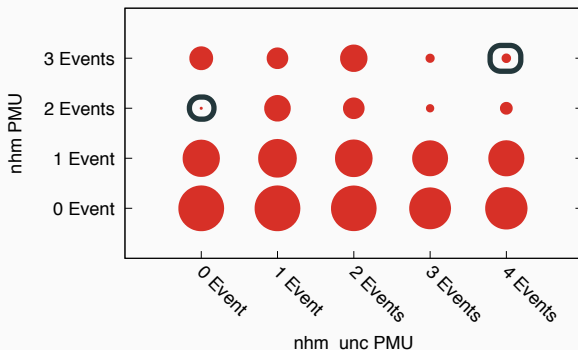




# OUR APPROACH: OPEN-TESTBED TO AUTOMATICALLY LEARN POWER MODELS

Average error per combination of HPCs for **freqmine**, **fluidanimate**, **facesim** on a Xeon W3520.


$$P_{idle} = 92\text{ W} ; P_{CPU} = \frac{1.40 \cdot \text{HPC (l1i:reads)}}{10^8} + \frac{7.29 \cdot \text{HPC (lsd:inactive)}}{10^9}$$



# OUR APPROACH: OPEN-TESTBED TO AUTOMATICALLY LEARN POWER MODELS

Relative errors for the PARSEC suite on the Cortex A15.

$$P_{idle} = 3.5 \text{ W} ; P_{CPU} = \frac{1.18 \cdot \text{cpu\_cycles}}{10^9} + \frac{1.26 \cdot \text{inst\_spec\_exec\_integer\_inst}}{10^{10}} + \frac{1.84 \cdot \text{bus\_cycles}}{10^{11}}$$



- **Portability**

Beyond SotA: 4 CPUs (2×Intel, 1 AMD, 1 **ARM**)

- **Portability**

Beyond SotA: 4 CPUs (2×Intel, 1 AMD, 1 ARM)

- **Accuracy**

Avg. error on the 4 CPUs: 1.5%

- **Portability**

Beyond SotA: 4 CPUs (2×Intel, 1 AMD, 1 ARM)

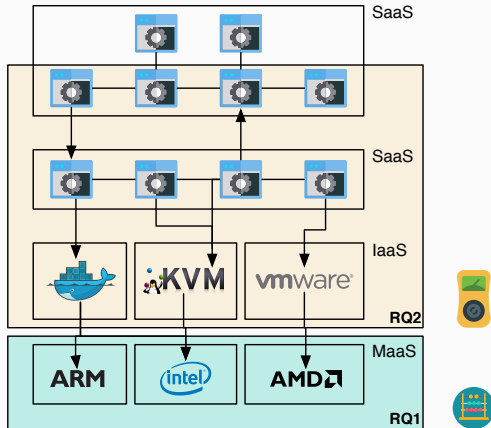
- **Accuracy**

Avg. error on the 4 CPUs: 1.5%


- **Reproducibility**

Built on open-source workloads

**RQ2:** Can we propose a uniform view of the service power consumption?




**RQ2:** Can we propose a uniform view of the service power consumption?



## Challenges

1. Native
2. Virtualized
3. Distributed

**RQ2:** Can we propose a uniform view of the service power consumption?



## Challenges

1. **Native**
2. Virtualized
3. Distributed



- Code freely available on GITHUB: <http://powerapi.org>
  - Scala / Akka
  - LoC: 8.7k
  - Docker
  - AGPLv3

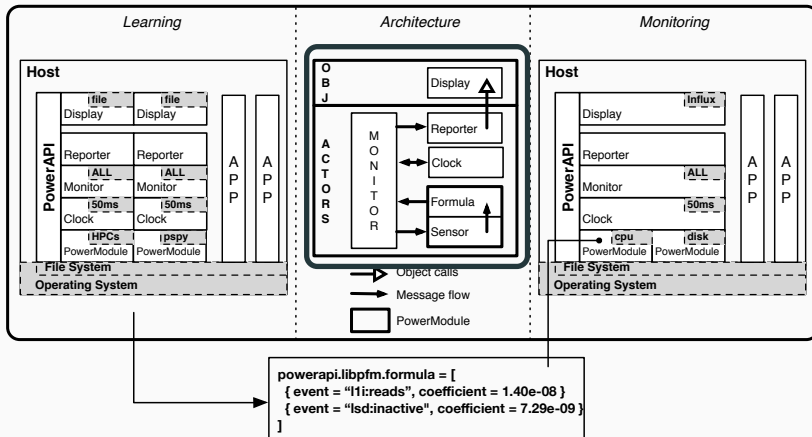
- Code freely available on GITHUB: <http://powerapi.org>
  - Scala / Akka
  - LoC: 8.7k
  - Docker
  - AGPLv3
- 2<sup>nd</sup> major iteration<sup>6</sup>
  - Full support of multi-core CPU architectures (HT, DVFS, TB)
  - Learning techniques
  - Better support of Akka

---

<sup>6</sup>A. Nouredine. “Towards a Better Understanding of the Energy Consumption of Software Systems”. PhD thesis. Université des Sciences et Technologie de Lille - Lille I, 2014.

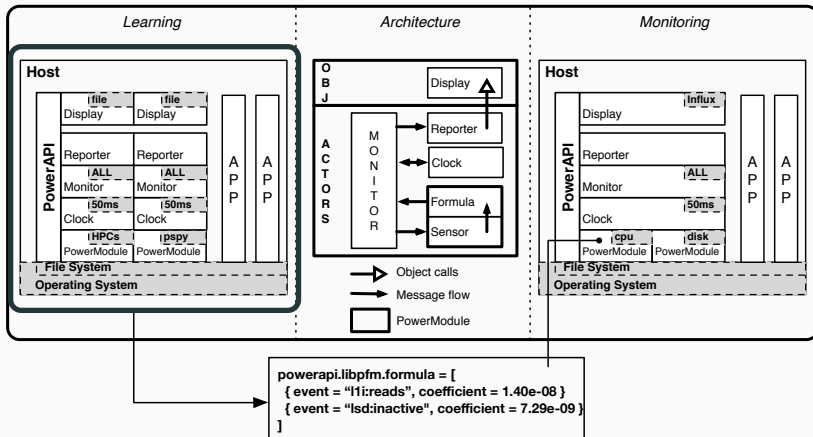
# POWERAPI: A TOOLKIT FOR BUILDING SD POWER METERS

POWERAPI's architecture & deployment.



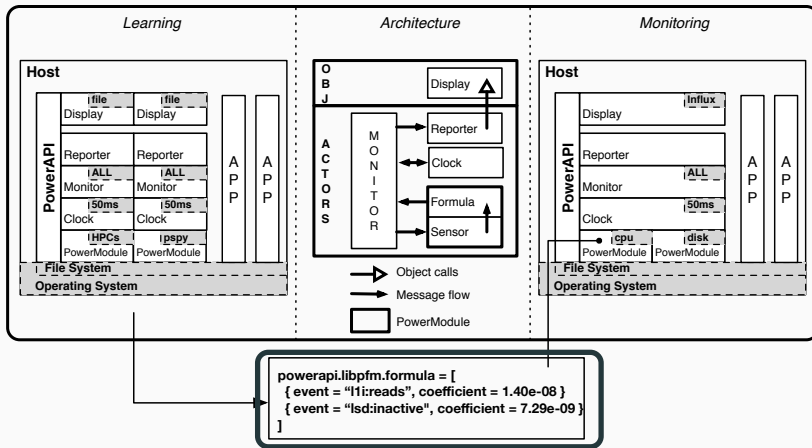
# POWERAPI: A TOOLKIT FOR BUILDING SD POWER METERS

POWERAPI's architecture & deployment.



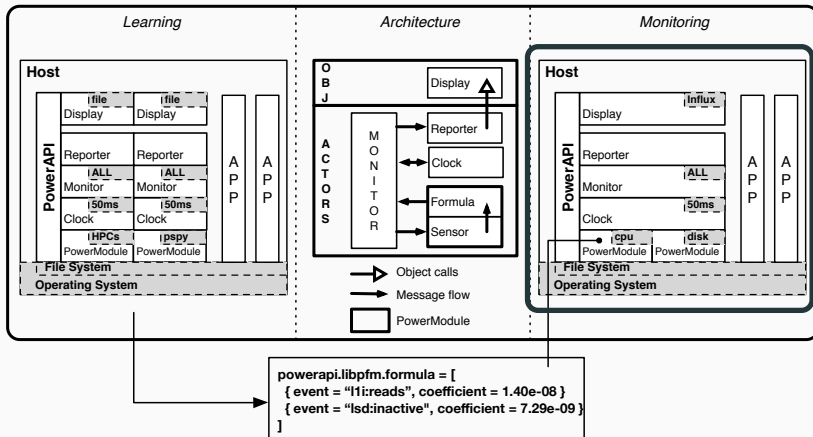
# POWERAPI: A TOOLKIT FOR BUILDING SD POWER METERS

## POWERAPI's architecture & deployment.




# POWERAPI: A TOOLKIT FOR BUILDING SD POWER METERS

POWERAPI's architecture & deployment.




# SD POWER METER FOR MONITORING CONCURRENT APPS



- On the Intel Xeon W3520
  - Monitoring freq.: 4Hz
  - Avg. error: 2%
  - Low overhead: 2 W

**RQ2:** Can we propose a uniform view of the service power consumption?




## Challenges


1. Native
2. **Virtualized**
3. Distributed



# BITWATTS ARCHITECTURE




Scaling PARSEC on multiple VMs on a Xeon W3520.



- Errors: from 1% (fluidanimate) up to 10% (swaptions)
- Beyond SotA [Ber+12]:<sup>7</sup> VM as a White-Box (+ multi-tenant)

<sup>7</sup>R. Bertran et al. "Energy Accounting for Shared Virtualized Environments Under DVFS Using PMC-based Power Models". In: *Future Generation Computer Systems* (2012).

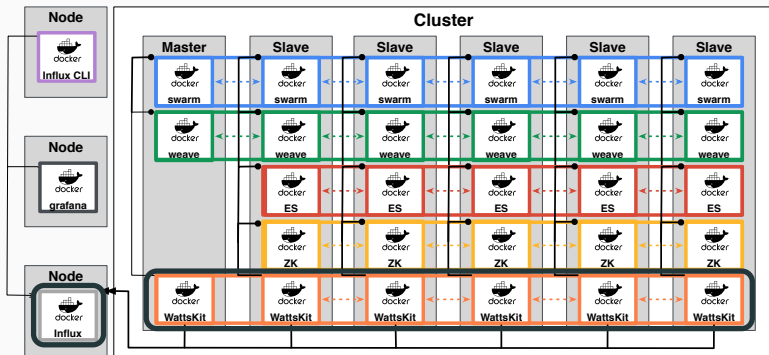
**RQ2:** Can we propose a uniform view of the service power consumption?



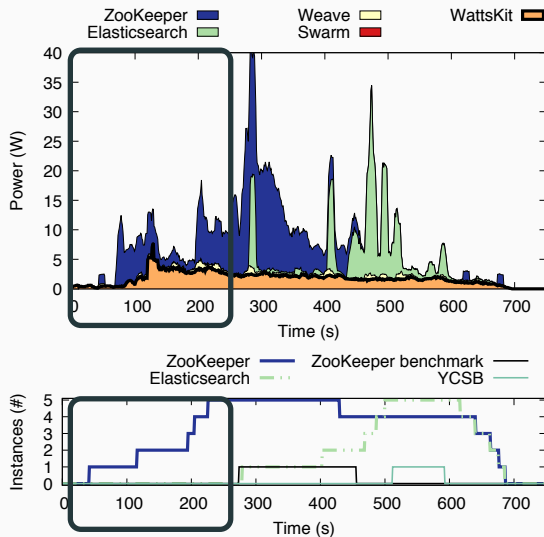
## Challenges

1. Native
2. Virtualized
3. **Distributed**


# A SERVICE-LEVEL POWER MONITORING




# A SERVICE-LEVEL POWER MONITORING



# A SERVICE-LEVEL POWER MONITORING



# A SERVICE-LEVEL POWER MONITORING



## CONCLUSION

---



POWERAPI, Deploying Software-defined Power Meters for the Cloud

POWERAPI, Deploying Software-defined Power Meters for the Cloud

- **RQ1:** Can we model the software power consumption regardless of the underlying architecture?

**Open-testbed approach for learning multi-core power models**

## POWERAPI, Deploying Software-defined Power Meters for the Cloud

- **RQ1:** Can we model the software power consumption regardless of the underlying architecture?

**Open-testbed approach for learning multi-core power models**

- **RQ2:** Can we propose a uniform view of the service power consumption?

**In width energy monitoring with POWERAPI, BITWATTS & WATTSKIT**



---

Maxime COLMANT — `maxime.colmant@inria.fr`

`http://powerapi.org`

---

- [Col+15] M. Colmant et al. “Process-level Power Estimation in VM-based Systems”. In: *Proceedings of the 10th European Conference on Computer Systems (EuroSys)*. 2015.
- [Hav+17] A. Havet et al. “GENPACK: A Generational Scheduler for Cloud Data Centers”. In: *IEEE International Conference on Cloud Engineering (IC2E)*. 2017.
- [Col+17b] M. Colmant et al. “WattsKit: Software-Defined Power Monitoring of Distributed Systems”. In: *17th IEEE/ACM International Symposium on Cluster, Cloud and Grid Computing (CCGrid)*. 2017.
- [Col+17a] M. Colmant et al. “The Next 700 CPU Power Models”. In: *ACM Trans. Model. Perform. Eval. Comput. Syst. (ACM TOMPECS)* (2017).



# REFERENCES I

- [Ber+10] R. Bertran et al. “Decomposable and Responsive Power Models for Multicore Processors Using Performance Counters”. In: *Proceedings of the 24th ACM International Conference on Supercomputing*. 2010.
- [Ber+12] R. Bertran et al. “Energy Accounting for Shared Virtualized Environments Under DVFS Using PMC-based Power Models”. In: *Future Generation Computer Systems* (2012).
- [BL09] C. Bienia and K. Li. “PARSEC 2.0: A New Benchmark Suite for Chip-Multiprocessors”. In: *Proceedings of the 5th Annual Workshop on Modeling, Benchmarking and Simulation*. 2009.
- [CM05] G. Contreras and M. Martonosi. “Power Prediction for Intel XScale® Processors Using Performance Monitoring Unit Events”. In: *Proceedings of the International Symposium on Low Power Electronics and Design*. 2005.
- [Col+15] M. Colmant et al. “Process-level Power Estimation in VM-based Systems”. In: *Proceedings of the 10th European Conference on Computer Systems (EuroSys)*. 2015.
- [Col+17a] M. Colmant et al. “The Next 700 CPU Power Models”. In: *ACM Trans. Model. Perform. Eval. Comput. Syst. (ACM TOMPECS)* (2017).
- [Col+17b] M. Colmant et al. “WattsKit: Software-Defined Power Monitoring of Distributed Systems”. In: *17th IEEE/ACM International Symposium on Cluster, Cloud and Grid Computing (CCGrid)*. 2017.

## REFERENCES II

- [CRS17] M. Colmant, R. Rouvoy, and L. Seinturier. “codEnergy: an Approach For Leveraging Source-Code Level Energy Analysis”. In: *To be chosen*. 2017.
- [Dol+15] M. F. Dolz et al. “An analytical methodology to derive power models based on hardware and software metrics”. In: *Computer Science - Research and Development* (2015).
- [ERK06] D. Economou, S. Rivoire, and C. Kozyrakis. “Full-System Power Analysis and Modeling for Server Environments”. In: *In Workshop on Modeling Benchmarking and Simulation*. 2006.
- [Hav+17] A. Havet et al. “GENPACK: A Generational Scheduler for Cloud Data Centers”. In: *IEEE International Conference on Cloud Engineering (IC2E)*. 2017.
- [IM03] C. Isci and M. Martonosi. “Runtime Power Monitoring in High-End Processors: Methodology and Empirical Data”. In: *Proceedings of the 36th Annual IEEE/ACM International Symposium on Microarchitecture*. 2003.
- [Int15] Intel. *Intel 64 and IA-32 Architectures Software Developer’s Manual*. 2015. URL: <http://www.intel.com/content/dam/www/public/us/en/documents/manuals/64-ia-32-architectures-software-developer-vol-1-manual.pdf> (visited on 08/01/2016).
- [Nou14] A. Nouredine. “Towards a Better Understanding of the Energy Consumption of Software Systems”. PhD thesis. Université des Sciences et Technologie de Lille - Lille I, 2014.

## REFERENCES III

- [RRK08] S. Rivoire, P. Ranganathan, and C. Kozyrakis. “A Comparison of High-level Full-system Power Models”. In: *Proceedings of the Conference on Power Aware Computing and Systems*. 2008.
- [The08] The Climate Group. *SMART 2020: Enabling the low carbon economy in the information age*. 2008. URL: <http://gesi.org/article/43> (visited on 09/23/2016).
- [Yan+14] H. Yang et al. “iMeter: An integrated VM power model based on performance profiling”. In: *Future Generation Computer Systems* (2014).
- [Zha+14] Y. Zhai et al. “HaPPy: Hyperthread-aware Power Profiling Dynamically”. In: *Proceedings of the USENIX Annual Technical Conference*. 2014.